

WPI

Human Context Recognition: A Controllable GAN Approach

Joshua DeOliveira, Worcester Polytechnic Institute

Harrison Kim, Northeastern University

MaryClare Martin, College of the Holy Cross

Faculty Advisor: Prof. Elke Rundensteiner

Ph.D. Mentor: Walter Gerych

Funded by NSF Grant #1852498

Progress Summary


1. Explored the ExtraSensory dataset through visualizations.
2. Further developed an understanding of PyTorch by building a simple NN classifier and a GAN for digit image generation.
3. Trained a Vanilla and Wasserstein GAN to generate accelerometer data in a non-sequential fashion.
4. Explored a novel training methodology for simple GANs.
5. Read and analyzed the strengths and weaknesses of different GAN-based architectures for time-series data.

ExtraSensory Dataset

- 60 Participants
- 300k+ data points
- 15 different devices
- 10 unique sensors
- 183 features
- 41 discrete device attributes
- 51 labels (not all mutually exclusive)



Accessible at: <http://extrasensory.ucsd.edu/>



Vaizman, Y., Ellis, K., and Lanckriet, G. "Recognizing Detailed Human Context In-the-Wild from Smartphones and Smartwatches." *IEEE Pervasive Computing*, October-December 2017.

Simple NN for Classification

Using featurized accelerometer data without temporal context from the ExtraSensory dataset, can we classify when users are sitting?

Model:

- 26 accelerometer features
- 4 ReLu layers (40 neurons each)
- 20% dropout
- 1 sigmoid layer to 1 neuron
 - Probability user was sitting

Training:

- $e = 0.001$
- Epochs = 120
- Training Loss: MSE
- Batch Size = 10,000 datum
- Volume = 264,142 datum

Results:

Naive Accuracy: 63.86%

Test Accuracy: 70.72% (+10.7%)

Loss After Training: 0.1882


PyTorch GAN for Digit Image Generation

Generator:

- Linear Layers
- 1D Batch Normalization
- ReLU Layers
- Sigmoid Output Layer

Discriminator:

- Linear Layers
- Leaky ReLU Layers
- Linear Output Layer
- BCE with Logits Loss



A Dynamic Training Approach

Reasoning:

Conventionally, GANs train by sequentially freezing one machine, **G** or **D**, and training the other, each at a constant number of epochs g and d . Ideally, $d > g$ and d trains first such that **D** is always better than **G** to provide effective feedback. This can lead to scenarios, after **G** begins to generate semi-believable data and **D** well understands the real data, keeping $d > g$ inefficiently wastes epochs that can lead to overfitting **D**. We can efficiently train both machine's by enforcing some reason of "fair competition" that requires **D** only barely better than **G** to maximize improving **G** (the purpose of our GAN) as many epochs as possible.

A Dynamic Training Approach

Hyper-parameters:

- Static_threshold: duration of Phase 1 (in epochs)
- D_static, G_static: # of epochs to train each model during static phase
- Pull_threshold: minimum fp Rate to keep **D** training
- Push_threshold: maximum fp Rate to keep **G** training
- Recall_threshold: minimum recall of **D** to begin Phase 3

A Dynamic Training Approach

Training Methodology:

Phase 1: Static Training

- Train conventionally to allow fundamental training to both machines based off of D_{static} and G_{static}

Phase 2: Check for Understanding

- Allow D to train until the recall of D is above the $Recall_threshold$

Phase 3: Pull / Push

1. Allow G to train until it fools D above the push threshold, then “push” G away
2. Allow D to train until it no longer is fooled by G according to the pull_threshold

Vanilla GAN for Generating non-Sequential Accelerometer Data

Generator:

- 4 ReLU/Batch Normalization layers
- 10% dropout
- Linear output layer (26-dimensional)

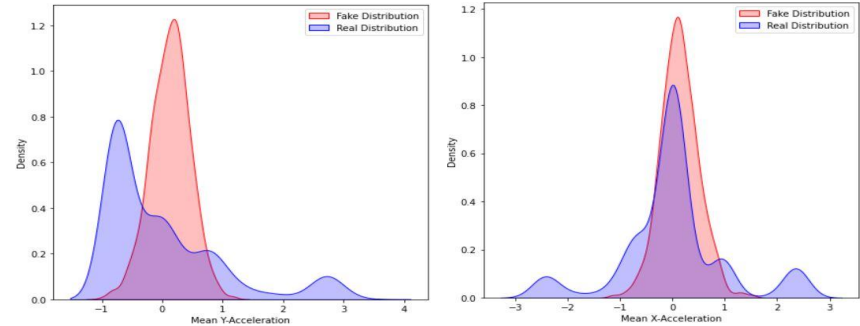
Discriminator:

- 3 Leaky ReLU layers
- 10% dropout
- Linear output layer (1-dimensional) with Sigmoid activation function
 - Probability of real/fake features

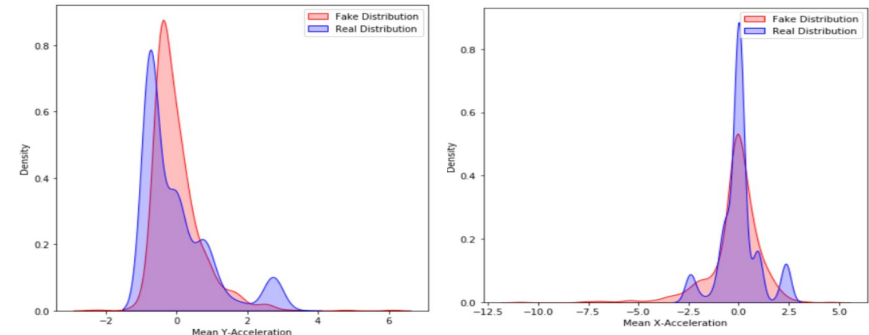
Training:

- Generator $e = 0.001$
- Discriminator $e = 0.0001$
- Epochs = 2000
- Training Loss = BCE Loss
- Batch Size = 1000 datum
- Latent Vector Dimension = 100



Before Training



2000 Epochs



Wasserstein Distance on Vanilla GAN



Wasserstein GAN for Generating non-Sequential Accelerometer Data

Generator:

- 4 ReLU/Batch Normalization layers
- 10% dropout
- Linear output layer (26-dimensional)

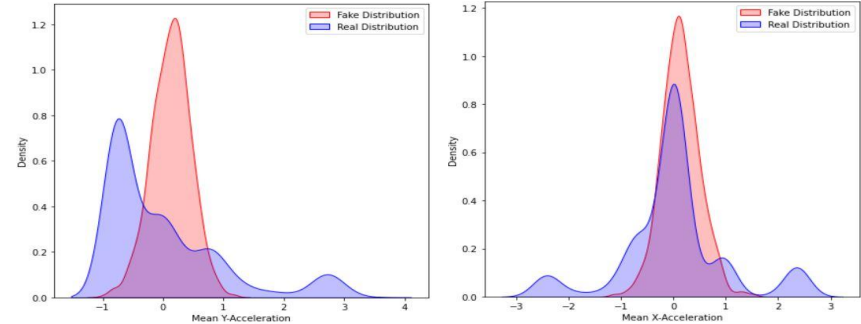
Discriminator:

- 3 Leaky ReLU layers
- 10% dropout
- Linear output layer (1-dimensional)

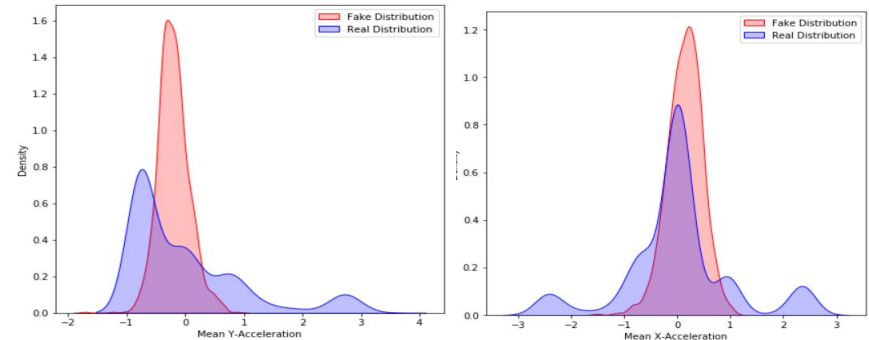
Training:

- Generator $\epsilon = 0.00001$
- Discriminator $\epsilon = 0.00001$
- Epochs = 800
- Training Loss = Wasserstein Distance
- Batch Size = 1000 datum
- Latent Vector Dimension = 100

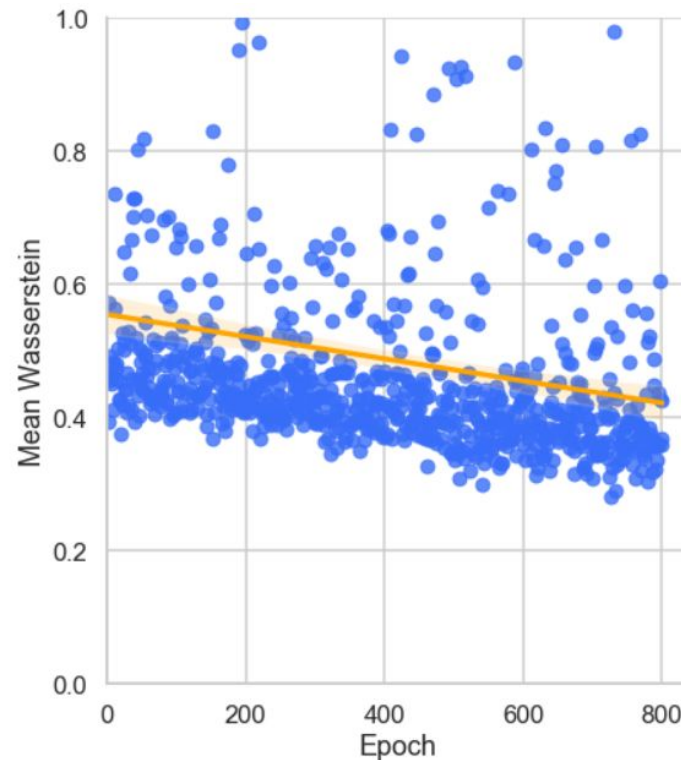

Before Training



800 Epochs




Wasserstein Distance on WGAN



GANs for Sequential Data

Generators with recurrent architectures allow for developing realistic time-series data for an iteration t by using previous iterations for generation ($t-1$, $t-2$, $t-3$, ... $t-n$)

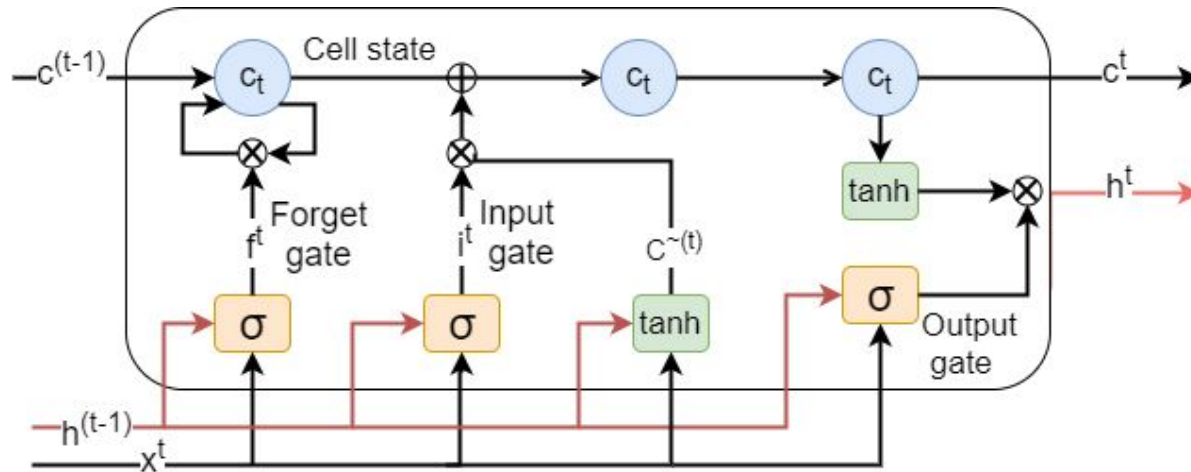
G



Feng et. al. (2017). Audio visual speech recognition with multimodal recurrent neural networks. 681-688. 10.1109/IJCNN.2017.7965918.

GANs for Sequential Data

Generators with recurrent architectures allow for developing realistic time-series data for an iteration t by using previous iterations for generation ($t-1$, $t-2$, $t-3$, ... $t-n$)



Jenkins et. al. (2018). Accident Scenario Generation with Recurrent Neural Networks. 3340-3345. 10.1109/ITSC.2018.8569661.

Next Steps

1. Implement third GAN evaluation metric (compare accuracy of a classifier trained on fake & real features to a classifier trained on only real features).
2. Continue hyperparameter tuning/experimentation of the Vanilla/Wasserstein GAN architectures.
3. Begin developing a Controllable GAN for generating accelerometer data for select classes.